

LEARNING RESOURCE GUIDE

ISC2 CCSP: Certified Cloud Security Professional (Aug 2026) | Professional Participant Workbook

Learning Objectives

By the end of this lecture and its accompanying guide, you will be able to:

- By the end of this lecture you will be able to distinguish Business Continuity from Disaster Recovery and explain the relationship between a BCP and a DRP.
- By the end of this lecture you will be able to describe the Business Impact Analysis process and explain why BIA outputs drive all BC/DR strategy decisions.
- By the end of this lecture you will be able to compare the traditional DR site types (hot, warm, cold) with cloud-native BC/DR patterns (backup and restore, pilot light, warm standby, active-active).
- By the end of this lecture you will be able to evaluate BC/DR strategies against cost and recovery speed trade-offs for a given organisational scenario.
- By the end of this lecture you will be able to explain the role of testing — tabletop through full interruption — in validating a BC/DR plan.
- By the end of this lecture you will be able to identify the BC/DR implications of multi-region, multi-availability-zone, and multi-cloud architectures.

Introduction

Business continuity and disaster recovery planning is the discipline that converts risk awareness into operational readiness. An organisation may understand that a regional cloud outage or a ransomware attack could disrupt its operations — but without documented, tested, and maintained BC/DR plans, that awareness provides no protection when a real incident occurs. Cloud computing has simultaneously increased the resilience options available to organisations and introduced new dependencies that make BC/DR planning more complex.

This lecture examines BC/DR strategy at the conceptual and architectural level: what BC and DR mean individually and together, how Business Impact Analysis generates the requirements that drive strategy selection, how the traditional hot/warm/cold site model maps to cloud-native patterns from backup and restore through active-active, and why testing is as operationally important as designing. This conceptual foundation is directly tested in Domain 3 of the CCSP examination and underpins the technical recovery objective content in the next lecture.

The practical stakes are high. The 2021 Kaseya ransomware attack disrupted over 1,500 organisations simultaneously. Organisations that had documented and tested BC/DR plans — including defined activation criteria, tested recovery procedures, and verified backup integrity — restored operations in hours to days. Organisations without those foundations faced weeks of downtime. The difference was not the quality of their security perimeter; it was whether they had invested in BC/DR as a practised capability rather than a document.

Core Concepts Explained

Business Continuity versus Disaster Recovery

Definition: Business Continuity is the broad discipline of ensuring that critical business functions remain operational before, during, and after a disruptive event. Disaster Recovery is a subset focused specifically on restoring IT systems and data after a disaster to re-enable business function.

Business relevance: The distinction matters because BC/DR investments must cover more than IT restoration. A business continuity plan may specify that customer service operations move to a third-party call centre during an outage, that paper-based processes substitute for digital ones in the short term, and that key suppliers are notified within defined timeframes — none of which are IT recovery steps, but all of which are essential to business survival.

Cloud Security Practitioner relevance: A Cloud Security Practitioner contributes primarily to the Disaster Recovery dimension — designing and validating the technical recovery architecture, ensuring cloud resources can be restored to meet recovery objectives, and verifying that recovery procedures are documented and executable. They must also communicate IT recovery capabilities and constraints to business continuity planners so that BCP assumptions align with what DRP can actually deliver.

Real-world example: During the 2017 AWS S3 US-East-1 service disruption — caused by a typo during a maintenance operation — many organisations discovered that their BCP assumed AWS S3 would be continuously available and had no documented process for handling S3 unavailability. Their DR plans covered server failures but not a regional storage service outage. Organisations that had separated their BC thinking (what does the business do if IT systems are unavailable?) from their DR thinking (how does IT restore systems?) were better prepared — they had pre-identified manual workarounds for critical business processes and communicated them to operations teams before the outage occurred.

Risks or limitations: Treating BC and DR as interchangeable leads to plans that cover IT recovery but fail to address what business units should do while IT is recovering. The best practice is to develop the BCP and DRP jointly but maintain them as separate documents with clearly delineated scope — the BCP owned by business leadership, the DRP owned by IT and security.

Business Impact Analysis

Definition: A structured process that identifies the organisation's critical business processes, quantifies the financial and operational impact of disruption to each process over time, and establishes the recovery time and data loss tolerances that BC/DR strategies must achieve.

Business relevance: The BIA is the foundation of BC/DR because it translates business risk into technical requirements. Without BIA outputs, IT teams cannot make defensible decisions about how much to invest in resilience for any given system. The BIA also produces the prioritised recovery order — which systems are restored first when resources are constrained during a real disaster.

Cloud Security Practitioner relevance: A Cloud Security Practitioner participates in the BIA by providing technical input on the cloud services that support each business process, identifying cloud-specific dependencies (cross-service, cross-account, cross-region), and translating business impact tolerances into technical specifications for backup frequency, replication configuration, and failover architecture.

Real-world example: A UK-based insurance company conducted a BIA before migrating its claims processing platform to AWS. Business stakeholders determined that claims processing being unavailable for more than two hours would trigger regulatory notification obligations under FCA rules, and that each additional hour of downtime cost approximately £200,000 in claims processing backlog and customer compensation. The BIA

output — a two-hour maximum tolerable downtime for claims processing — directly drove the technical decision to implement active-active multi-region deployment for that workload. The BIA for the internal HR system, by contrast, showed a 72-hour maximum tolerable downtime with negligible financial impact, justifying a daily backup-and-restore approach. Without the BIA, the cloud security team would have had no objective basis for allocating resilience investment across these two very different workloads.

Risks or limitations: BIAs based on IT assumptions rather than business stakeholder input often underestimate the impact of disruption to customer-facing processes and overestimate tolerance for downtime in back-office systems. The best practice is to conduct BIA interviews with business unit owners, not IT managers, and to validate financial impact figures with the finance function before using them to justify resilience investment.

Traditional DR Site Types

Definition: The classification of disaster recovery facilities into hot sites (fully operational duplicates with near-real-time replication, enabling failover in minutes), warm sites (hardware and software installed but not running, enabling failover in hours), and cold sites (facilities with connectivity but no hardware, requiring days to weeks to stand up).

Business relevance: The hot/warm/cold framework establishes the fundamental trade-off between recovery speed and cost that governs all DR architecture decisions. Understanding this framework is necessary to understand how cloud-native patterns improve on it.

Cloud Security Practitioner relevance: A Cloud Security Practitioner must be able to map organisational recovery requirements to DR site type equivalents and then identify the cloud-native implementation pattern that achieves the required recovery speed at the lowest sustainable cost.

Real-world example: Before migrating to cloud, a healthcare organisation maintained a hot site forty miles from its primary data centre, running synchronously replicated copies of its electronic health record system on leased hardware. The cost was £1.2 million per year in facility fees, hardware maintenance, and network connectivity — even though the site was used for actual DR purposes only once in seven years. After migration to AWS, the organisation replaced the hot site with an active-active multi-region deployment using Aurora Global Database for synchronous replication. The annual cost dropped to approximately £180,000 in additional AWS compute and data transfer charges while achieving the same recovery time. The traditional hot site model's cost was driven by dedicated physical infrastructure that sat idle; cloud's cost model charges only for resources consumed.

Risks or limitations: Cold sites carry the risk that hardware procurement and software installation during an actual disaster will take longer than the documented estimate, especially if the disaster is regional and affects supply chains. The best practice for organisations still operating cold sites is to conduct an annual full-scale test of standing up the cold site and document the actual time taken — not the theoretical estimate.

Cloud-Native BC/DR Patterns

Definition: The spectrum of cloud-based disaster recovery architectures, from backup and restore (lowest cost, longest recovery time) through pilot light (core components always running) and warm standby (reduced-capacity environment always running) to active-active (full capacity in multiple locations simultaneously).

Business relevance: Cloud-native patterns decouple the cost-recovery-speed trade-off that constrains traditional DR site decisions. By using on-demand infrastructure provisioning and infrastructure as code, organisations can achieve warm or hot site recovery speeds at costs approaching cold site levels for the backup and restore pattern, or achieve near-zero recovery time with active-active at costs that scale with actual capacity rather than dedicated idle hardware.

Cloud Security Practitioner relevance: A Cloud Security Practitioner must understand the security implications of each pattern: backup and restore requires backup encryption and access control; pilot light requires that the always-on components (database replicas, network configurations) are kept current and secured; warm standby requires continuous monitoring of the reduced-capacity environment; active-active requires consistent security controls across all active sites and careful management of data synchronisation integrity.

Real-world example: A retail company ran a pilot light BC/DR architecture on AWS for its e-commerce platform. The pilot light consisted of: an RDS Aurora read replica in eu-west-2 (London) replicating from the primary in eu-west-1 (Ireland), an AMI of the application tier updated weekly via automated build pipeline, and pre-configured VPC, IAM, and Route 53 configurations held in Terraform templates. In January 2023, a major fire in the eu-west-1 data centre used by their on-premises networking provider disrupted connectivity to their AWS environment — not an AWS failure, but an effectively equivalent impact. The team promoted the Aurora read replica to primary, ran the Terraform templates to provision the application tier in eu-west-2, and updated Route 53 to point traffic to the London environment. Total recovery time was two hours and twelve minutes, compared to their RTO of four hours. The test confirmed that the pilot light strategy met their business requirements at a cost approximately 60% lower than warm standby would have required.

Risks or limitations: The backup and restore pattern's recovery time depends heavily on data volume. As data grows, restore time may silently exceed RTO unless regularly tested. The best practice is to run quarterly restore tests that simulate the actual data volume expected at recovery time and verify that the restore process completes within the RTO.

BC/DR Plan Testing

Definition: The structured programme of exercises — from tabletop through full interruption — that validates whether a BC/DR plan will achieve its objectives under realistic conditions.

Business relevance: An untested BC/DR plan is not a recovery capability — it is a document. Testing reveals gaps between planned and actual recovery times, identifies team members who do not understand their roles, surfaces dependencies that the plan does not cover, and builds the muscle memory that enables effective response under the stress of an actual incident.

Cloud Security Practitioner relevance: A Cloud Security Practitioner must design and participate in BC/DR tests that validate cloud-specific recovery procedures — automated failover mechanisms, infrastructure-as-code deployment pipelines, database promotion procedures, and DNS cutover processes — not just the organisational communication steps that tabletop exercises typically focus on.

Real-world example: A financial technology firm conducted its first parallel BC/DR test for its cloud payment processing platform eighteen months after completing the DR architecture design. The test revealed three critical gaps: the Terraform deployment pipeline had not been updated to reflect a database schema change made six months earlier, so the recovery environment came up but the application could not connect to the database; the DNS TTL for the payment gateway endpoint was set to 3,600 seconds (one hour) rather than the assumed 60 seconds, meaning payment traffic would not redirect for a full hour after failover; and the team member responsible for the database promotion step had left the company and no one else had been trained on the procedure. None of these gaps would have been visible in a tabletop exercise. The parallel test cost approximately one additional day of engineering time but potentially saved days of extended downtime in a real disaster. The best practice demonstrated here is that parallel tests must exercise the actual technical runbooks, not just the decision-making process.

Risks or limitations: Organisations often default to tabletop exercises for all BC/DR testing because they are low-cost and low-risk. Tabletop exercises are valuable but insufficient — they cannot validate that automated recovery mechanisms work, that runbook commands are accurate, or that recovery time estimates based on

current data volumes are still achievable. The best practice is to run at least one parallel test annually that activates actual recovery infrastructure.

Practical Applications

The following scenarios illustrate how Cloud Security Practitioners apply BC/DR strategy principles in real organisational contexts.

Selecting a Cloud BC/DR Pattern Based on BIA Outputs

A Cloud Security Practitioner at a logistics company is asked to design the DR architecture for three workloads migrating to AWS: a real-time shipment tracking API used by enterprise customers (BIA output: RTO 30 minutes, financial impact £50,000 per hour of downtime), a warehouse management system (BIA output: RTO 4 hours, impact £8,000 per hour), and a business intelligence reporting platform (BIA output: RTO 48 hours, impact negligible). For the shipment tracking API, the practitioner recommends active-active deployment across two AWS regions with Route 53 latency-based routing and Aurora Global Database — the 30-minute RTO cannot be met by any architecture that requires manual failover steps. For the warehouse management system, she recommends a warm standby in a second region with an RDS read replica promoted to primary in the event of a regional failure — the four-hour RTO provides enough time for a partially manual failover process. For the reporting platform, she recommends daily snapshot backups to S3 Cross-Region Replication with a documented restore runbook — the 48-hour RTO and minimal financial impact do not justify the cost of always-on standby infrastructure. The BIA outputs are the sole justification for each architectural choice, making the security investment defensible to finance and executive leadership.

Updating BC/DR Plans After Infrastructure Changes

A Cloud Security Practitioner at a media company is responsible for BC/DR plan maintenance. After a six-month platform modernisation project that replaced the company's legacy database tier with Amazon DynamoDB Global Tables, he schedules a BC/DR plan review. He discovers that the existing DRP references RDS read replica promotion as the database recovery procedure — a procedure that no longer applies to DynamoDB. He also discovers that the BIA used to set the RTO for the video streaming platform assumed a six-hour restore time for the database tier, based on the old RDS backup restore speed. DynamoDB Global Tables enable single-digit millisecond failover, meaning the database is no longer the constraint on recovery time — the application tier deployment is. He updates the DRP with the correct DynamoDB failover procedure (updating the application's regional endpoint configuration), revises the RTO to two hours based on application tier provisioning time, and schedules a tabletop exercise to validate the updated plan with the infrastructure team. This example illustrates the best practice that BC/DR plans must be updated as a required output of every significant infrastructure change — not reviewed on a fixed annual schedule regardless of change activity.

Designing Multi-Cloud BC/DR for Provider Resilience

A Cloud Security Practitioner at a government contractor is asked to evaluate a multi-cloud BC/DR architecture that would enable the organisation to fail over from AWS to Azure in the event of an extended AWS regional outage. She begins by documenting the dependencies: the application workloads are containerised and run on Kubernetes, making them relatively portable; the databases use PostgreSQL, which runs on both AWS RDS and Azure Database for PostgreSQL; but the organisation's identity architecture uses AWS IAM and AWS Cognito, which have no direct Azure equivalent. She also notes that data replication across cloud providers requires traversing the public internet or establishing dedicated cross-cloud connectivity, introducing latency and cost that same-provider cross-region replication does not have. Her recommendation is a hybrid approach: for

workloads where multi-cloud failover is technically feasible and the BIA justifies the cost, implement synchronous database replication to Azure with containerised application deployment validated quarterly. For IAM and identity, accept that multi-cloud failover would require re-establishing identity infrastructure in Azure and document that as a gap with a four-hour recovery allowance. The key best practice here is that multi-cloud BC/DR complexity must be weighed against actual risk — a provider-level outage affecting a whole region is a low-probability event, and the additional operational complexity of multi-cloud management must be justified by the business impact of that scenario.

Best Practices

- **Drive all BC/DR strategy decisions from BIA outputs, not from available budget or vendor recommendations.** The BIA establishes what each system's downtime actually costs the business. Use those figures to calculate the maximum defensible investment in resilience for each workload. A system whose downtime costs £1,000 per hour does not justify a £500,000 active-active architecture; a system whose downtime costs £100,000 per hour does. This discipline prevents both over-investment and under-investment and makes BC/DR spending auditable.
- **Use infrastructure as code to keep recovery environments synchronised with production environments.** The most common cause of BC/DR test failures is that the recovery environment has drifted from production — a configuration change, a schema migration, or a new service dependency was made in production but not reflected in the DR runbooks. Managing both environments from the same IaC repository and running the DR deployment pipeline in CI/CD reduces drift and ensures that the recovery environment is an accurate representation of what production needs.
- **Test at a level that exercises actual technical procedures, not just decision-making processes.** Tabletop exercises are valuable for testing communication flows, escalation paths, and decision authority. They are insufficient for validating recovery time estimates, runbook accuracy, or automated failover mechanisms. Conduct at least one parallel or full interruption test annually for each high-criticality workload, and treat test failures as bugs in the plan that must be fixed before the next test.
- **Document activation criteria explicitly, including who has authority to declare a disaster and invoke the plan.** Ambiguity about activation authority causes dangerous delay during actual incidents. BC/DR plans must specify: what conditions trigger activation (e.g. RTO threshold exceeded, confirmed ransomware encryption of production data, provider declares regional disaster), who can declare a disaster (named roles, not individuals), and what communication steps must happen within the first thirty minutes of activation.
- **Account for cloud provider dependencies in the BIA and plan for provider outage scenarios.** An organisation that has moved all workloads to a single cloud provider has created a concentration dependency that the BIA must quantify. Understand the provider's published availability SLAs, their historical incident record for the regions you use, and what the business impact would be if a regional outage lasted four hours, twelve hours, or twenty-four hours. Use those scenarios to validate that your chosen BC/DR pattern provides adequate resilience for the provider dependency risk.

Common Mistakes and Misconceptions

Misconception: Selecting a BC/DR strategy (e.g. backup and restore) before conducting a BIA, and then discovering that the strategy cannot meet the organisation's actual recovery requirements.

Correct approach: Always conduct the BIA first. Recovery strategy selection should be a consequence of BIA outputs — specifically the RTO and RPO for each system — not a starting point. Starting with a budget or a

preferred architecture and then checking whether it meets requirements is a backwards process that consistently produces gaps.

Misconception: Assuming that cloud provider availability SLAs guarantee recovery within the organisation's RTO.

Correct approach: Provider SLAs define the availability commitment for the cloud service, not the time required to restore a specific customer's workload after a failure. A provider SLA of 99.99% uptime means the service is expected to be unavailable for less than 52 minutes per year — it does not mean customer workloads will be restored within any particular timeframe. Customer workload recovery times depend on the architecture the customer has built, not the provider's SLA.

Misconception: Failing to update BC/DR plans and recovery environments after infrastructure changes, resulting in runbooks that reference systems, procedures, or configurations that no longer exist.

Correct approach: Integrate BC/DR plan review into the change management process. Every significant infrastructure change — new service adoption, database migration, architecture refactoring — should trigger a review of the DRP runbooks and recovery environment configuration to verify that they remain accurate. Use infrastructure as code to reduce the risk of drift.

Misconception: Running only tabletop exercises and treating them as sufficient validation of BC/DR readiness.

Correct approach: Tabletop exercises test the decision-making and communication aspects of BC/DR but cannot validate that recovery procedures work technically. Conduct parallel tests that activate actual recovery infrastructure to confirm that automated failover mechanisms, runbook commands, and recovery time estimates are accurate.

table: Cloud BC/DR Pattern Comparison Matrix

Use this matrix to select the appropriate cloud BC/DR pattern for a given workload based on its BIA-defined recovery requirements and cost tolerance.

Pattern	Cloud Equivalent	Typical RTO	Typical RPO	Relative Cost	Best Suited For
Cold Site	Backup and Restore	Hours to days	Hours	Lowest	Non-critical internal systems; relaxed RTO/RPO
Warm Site	Pilot Light	1–4 hours	Minutes to 1 hour	Low-Medium	Important systems; moderate RTO tolerance
Hot Site (partial)	Warm Standby	Minutes to 1 hour	Near-zero to minutes	Medium-High	Business-critical systems; tight RTO
Hot Site (full)	Active-Active	Near-zero (seconds)	Near-zero	Highest	Mission-critical; zero-downtime requirement

How to use this: Cost relativities are approximate and depend on data volumes, compute sizes, and provider pricing. Always validate cost estimates with actual provider pricing before committing to an architecture.

Knowledge Check

Test your understanding before moving on. Answers are shown below each question.

Q1. A Cloud Security Practitioner's organisation operates a payment processing platform. The BIA team determines that every hour of downtime costs £150,000 in lost transactions and regulatory penalties. Which cloud BC/DR pattern is most likely justified by this finding?

- A) A. Backup and restore to a secondary region using daily snapshots
- B) B. Pilot light with a database replica and pre-built AMIs in a secondary region
- C) C. Active-active deployment across two regions with automated traffic failover
- D) D. Cold site equivalent with infrastructure-as-code templates ready to deploy

Answer: C. Active-active deployment across two regions with automated traffic failover

Q2. During a BC/DR tabletop exercise, the team walks through a scenario in which the primary AWS region is unavailable. A team member asks how long the database promotion step will take. No one in the room knows. What does this reveal about the BC/DR plan, and what should the team do?

- A) A. The plan is adequate — the team can look up the procedure in the runbook during a real incident
- B) B. The plan has a gap — the runbook must document the exact database promotion procedure with estimated execution time, and the team must validate it through a parallel test
- C) C. The plan is adequate — the cloud provider will handle database promotion automatically under the SLA
- D) D. The team should escalate to the cloud provider's support team, who will provide the procedure during a real incident

Answer: B. The plan has a gap — the runbook must document the exact database promotion procedure with estimated execution time, and the team must validate it through a parallel test

Q3. An organisation's DRP specifies that the recovery environment will be provisioned using Terraform templates. Six months after the DRP was written, the production environment was migrated from PostgreSQL 13 on RDS to Aurora PostgreSQL 15. Which risk does this create for the BC/DR plan?

- A) A. The Terraform templates reference the old database engine version and may fail to provision the correct recovery database
- B) B. Aurora PostgreSQL is not supported in the secondary region specified in the DRP
- C) C. PostgreSQL 15 cannot be restored from the backups taken under PostgreSQL 13
- D) D. The migration invalidated the provider's availability SLA for the primary region

Answer: A. The Terraform templates reference the old database engine version and may fail to provision the correct recovery database

Q4. What is the primary output of a Business Impact Analysis that directly drives BC/DR strategy selection?

- A) A. A list of the threats most likely to cause disruption to cloud infrastructure
- B) B. The financial and operational impact of disruption to each critical business process, expressed as maximum tolerable downtime and data loss tolerance
- C) C. A comparison of cloud provider availability SLAs against the organisation's uptime requirements
- D) D. The total annual cost of BC/DR infrastructure investments across all workloads

Answer: B. The financial and operational impact of disruption to each critical business process, expressed as maximum tolerable downtime and data loss tolerance

Q5. A Cloud Security Practitioner conducts a full interruption BC/DR test in which production is actually failed over to the recovery environment. The test reveals that recovery took six hours — two hours beyond the defined RTO of four hours. What should happen next?

- A) A. Revise the RTO upward to six hours to reflect actual recovery capability

- B) B. Document the result as acceptable variance since the test was stressful
- C) C. Conduct a root cause analysis of the recovery delay, update the plan and architecture to address the root cause, and retest before the next annual test cycle
- D) D. Escalate to the cloud provider's support team to reduce their platform recovery time

Answer: C. Conduct a root cause analysis of the recovery delay, update the plan and architecture to address the root cause, and retest before the next annual test cycle

Reflection Questions

1. Consider an organisation you know or work in. If its primary cloud region went offline for 24 hours today, which business processes would be most severely affected? Does that match what a formal BIA would say?
2. Think about the cloud BC/DR pattern your organisation currently uses (or would use). How was that pattern selected? Was it driven by BIA outputs, or was it selected for other reasons? What would a BIA-driven review reveal?
3. Reflect on the last significant infrastructure change in a cloud environment you know. Was the BC/DR plan reviewed and updated as part of that change? What might have been missed?
4. The lecture notes that testing is as important as designing. What is the most realistic BC/DR test your organisation currently runs? What gap would only be revealed by a more rigorous test?
5. Multi-cloud BC/DR is appealing in theory but operationally complex. Under what circumstances would you recommend it to a business stakeholder? What criteria would need to be met to justify the additional complexity?
6. How would you explain the difference between a cloud provider's availability SLA and your organisation's recovery capability to a non-technical executive who assumes that a 99.99% SLA means their workloads will be restored in under an hour?

Key Takeaways

- Business Continuity is the broad discipline of maintaining critical business functions during and after disruption; Disaster Recovery is the IT-focused subset that restores systems and data.
- The Business Impact Analysis is the mandatory starting point for BC/DR — it quantifies the financial and operational impact of disruption and establishes the RTO and RPO requirements that all strategy decisions must satisfy.
- Traditional DR site types (hot, warm, cold) map to cloud-native patterns (active-active, warm standby, pilot light, backup and restore) with the key advantage that cloud patterns decouple cost from recovery speed through on-demand provisioning.
- BC/DR strategy selection must be driven by BIA outputs — specifically the RTO and RPO for each system — not by available budget or vendor preferences.
- Cloud BC/DR introduces provider dependency as a BC/DR risk that must be explicitly addressed in the BIA and mitigated through multi-availability-zone, multi-region, or multi-cloud architectures proportional to business risk.
- Testing is a mandatory component of BC/DR readiness — tabletop exercises test decision-making; parallel and full interruption tests validate technical procedures and recovery time estimates.
- BC/DR plans must be treated as living documents reviewed and updated after every significant infrastructure change, not on a fixed annual schedule.

- Communication planning — who notifies whom, in what order, through what channels — is an essential and often overlooked element of BC/DR that must be documented and tested alongside technical recovery procedures.

Recommended Next Actions

Put this lecture to work within the next seven days:

7. Identify one business-critical cloud workload in your environment this week and verify whether a formal BIA has established its RTO and RPO — if not, initiate a conversation with the workload owner to start that process.
8. Review the BC/DR test schedule for at least one high-criticality system and verify whether the most recent test exercised actual technical recovery procedures (parallel or full interruption) or only tabletop decision-making — schedule an upgrade if needed.
9. Check whether the BC/DR runbooks for one workload have been updated since the last significant infrastructure change — identify any drift between the current architecture and what the runbooks describe.
10. Confirm that activation criteria and declaration authority for your most critical BC/DR plan are unambiguous and documented — verify that the named decision-makers are current employees and that their contact details are up to date.