

LEARNING RESOURCE GUIDE

ISC2 CCSP: Certified Cloud Security Professional (Aug 2026) | Professional Participant Workbook

Security Operations Center (SOC) and Intelligent Monitoring

Domain 5: Cloud Security Operations | ISC2 CCSP Exam Preparation

1. Learning Objectives

- By the end of this lecture, you will be able to describe the core mission and operational layers of a cloud Security Operations Center.
- By the end of this lecture, you will be able to explain how intelligent monitoring differs from signature-based detection and why both are necessary.
- By the end of this lecture, you will be able to identify the monitoring technologies used in a cloud SOC, including IDS, IPS, honeypots, and network security group logs.
- By the end of this lecture, you will be able to describe the role of machine learning and behavioural analytics in detecting novel or low-and-slow threats.
- By the end of this lecture, you will be able to explain why human judgment remains essential despite increased automation in SOC operations.

2. Introduction

At 02:43 on a Tuesday morning, an anomaly detection system at a global logistics company flags an unusual sequence of API calls. A service account that has made the same 12 API calls in the same order every night for eight months is making 47 different calls — including calls to data export endpoints it has never touched. The alert reaches the on-call SOC analyst at 02:44. By 02:51, the account is suspended and the incident ticket is open. The attacker had a seven-minute window. Without the SOC's continuous monitoring infrastructure, they would have had the entire night.

A Security Operations Center is the organisational capability responsible for continuous monitoring and response. In cloud environments, that capability spans multiple technical layers, combines automated detection with human analysis, and must operate without interruption across all time zones. The SOC is not a physical room, though it may have one. It is a function: a set of people, processes, and technologies working together to detect, triage, and respond to threats before they become incidents, and to contain incidents before they become catastrophes.

Understanding what a mature cloud SOC looks like, how its monitoring layers interact, and where intelligent monitoring changes the detection calculus is essential for cloud security professionals. This is also an area ISC2 tests directly, often by asking candidates to identify

which monitoring capability would detect a specific type of threat, or to evaluate whether a described SOC design is adequate for a given risk scenario.

3. Core Concepts Explained

3.1 The SOC's Core Mission: Continuous Monitoring and Response

The SOC's defining characteristic is continuous operation. Threats do not observe business hours, and attackers frequently time their activity for the periods of lowest expected human attention — weekends, holidays, and the quiet hours between midnight and 06:00 local time.

Business relevance: An organisation that monitors its cloud environment only during business hours has predictable blind spots. Sophisticated attackers know exactly where those blind spots are and time their most sensitive activities accordingly. Continuous monitoring is not a luxury feature — it is the baseline requirement for meaningful threat detection.

Practitioner angle: Continuous monitoring in a cloud environment requires automation for initial detection, triage workflows that function without the full team present, and on-call rotation procedures that ensure a qualified analyst is always reachable. The human layer does not need to be fully staffed at 03:00, but it must be reachable.

Risk: SOC teams that rely on individual analysts to catch threats through manual log review during business hours are not operating a SOC. They are operating a reactive log review function. The distinction matters because the two have fundamentally different detection latencies.

3.2 Monitoring Layers in a Cloud SOC

A mature cloud SOC monitors simultaneously across multiple technical layers. Each layer provides a different view of the environment, and effective detection requires correlating across all of them.

Firewalls and network security groups: In cloud platforms such as AWS and Azure, network security groups (NSGs) control traffic at the subnet and interface level and generate flow logs showing which connections were allowed or denied, source and destination addresses, ports, and protocol. These logs are the foundation of network-level visibility.

Intrusion Detection Systems and Intrusion Prevention Systems: IDS identifies and alerts on suspicious traffic patterns and known attack signatures. IPS goes further by actively blocking identified threats before they complete. In cloud environments, these functions are often delivered as cloud-native services or as virtual appliances deployed within the environment.

Honeypots: A honeypot is a deliberately exposed decoy system or resource with no legitimate business purpose. Any interaction with a honeypot is, by definition, suspicious. Honeypots generate high-fidelity alerts because there is no benign reason for a legitimate user or process to interact with them. Cloud environments make honeypots relatively easy to deploy — a storage bucket with a misleading name, an internal service endpoint that should never receive external traffic, or a set of credentials that exist only as bait.

Identity and access management logs: Authentication events, privilege escalation, and access pattern changes are among the highest-value signals for detecting account compromise and insider threats. Cloud IAM logs are essential and often underutilised.

3.3 Signature-Based Detection vs. Intelligent Monitoring

Traditional security monitoring relied on signature-based detection: incoming data is compared against a library of known bad patterns, and alerts fire when a match is found. Signature-based detection is fast, accurate for known threats, and generates low false positive rates against threats it knows about. Its fundamental limitation is that it cannot detect what it has not been told about.

Intelligent monitoring addresses this limitation. Using machine learning and behavioural analytics, it builds a statistical model of what normal looks like in a specific environment — the typical login times for each account, the standard API call patterns for each service, the normal volume and destination distribution of outbound network traffic. Once the baseline is established, deviations from it trigger alerts, regardless of whether the deviation matches a known attack signature.

Business relevance: Novel attacks, zero-day exploits, and low-and-slow intrusions — attacker activity spread across weeks to avoid triggering volume-based thresholds — do not match known signatures. They do, however, deviate from normal behaviour. An account that logs in from London every morning and appears in Singapore at 03:00 is a deviation. A server that has never made outbound connections to external IPs and suddenly starts doing so is a deviation. Intelligent monitoring catches what signature-based systems miss.

Real-world example: AWS GuardDuty uses machine learning to continuously analyse CloudTrail API call logs, VPC flow logs, and DNS logs. It identifies cryptocurrency mining behaviour, credential exfiltration patterns, and reconnaissance activities without requiring manual rule configuration for each threat type. This kind of automated, AI-driven monitoring is becoming a standard expectation rather than an advanced feature, and cloud security professionals should understand both its capabilities and its limitations.

3.4 The Role of Human Analysts

The most common misconception about intelligent monitoring is that it eliminates the need for human analysts. It does not. Intelligent monitoring reduces alert volume and improves signal quality by filtering noise. It does not eliminate false positives entirely, and it cannot make the judgment calls that matter most.

Human analysts are responsible for triage — evaluating whether a specific anomaly represents a genuine threat or a legitimate business change that the system has not yet learned to recognise. They handle escalation decisions, coordinate containment actions, and make the contextual judgments about whether a detected behaviour is malicious or explainable. These are not tasks that can be automated away. The value of intelligent monitoring is that it ensures human analysts spend their time on anomalies that warrant attention, not on obvious false positives.

4. SOC Monitoring Layer Comparison

Monitoring Layer	What It Detects	Cloud Implementation	Key Limitation
Signature-based IDS/IPS	Known attack patterns and malware signatures	Cloud-native services or virtual appliances	Cannot detect novel or zero-day threats
Behavioural Analytics / ML	Deviations from established normal behaviour	AWS GuardDuty, Azure Sentinel ML rules	Requires baseline period; generates false positives during legitimate change
Network Security Group Logs	Allowed/denied connections, traffic patterns	AWS VPC Flow Logs, Azure NSG flow logs	Volume — requires filtering and correlation to be actionable
IAM / Authentication Logs	Account compromise, privilege escalation, unusual access times	CloudTrail, Azure Activity Log, identity provider logs	Context-dependent; legitimate changes can mimic suspicious patterns
Honeypots	Any unauthorised interaction or reconnaissance	Decoy buckets, fake endpoints, canary credentials	Narrow scope — only covers resources specifically configured as decoys
Endpoint Telemetry	Host-level process activity, file changes, network connections	Cloud-hosted agents, EDR solutions	Coverage gaps where agents cannot be deployed

5. Case Study

Case Study: A major cloud provider's internal threat intelligence team documented a series of compromises at mid-market organisations in which attackers used credential stuffing to gain initial access, then spent between three and six weeks performing low-volume reconnaissance before exfiltrating data. In each case, the activity level was carefully kept below volume-based alert

thresholds. None of the affected organisations' signature-based IDS systems generated alerts. The organisations that detected the intrusion during the reconnaissance phase, before exfiltration began, shared one common characteristic: all had deployed behavioural analytics with established baselines. The anomaly flagged in each case was not the volume of activity — it was the type. A service account that had never accessed the data warehouse began running queries against it every night at 01:00. No signature matched. The behavioural model flagged a deviation from established access patterns. The lesson: volume-based thresholds alone are insufficient for detecting patient, deliberate attackers. Behavioural baselines catch what thresholds miss.

6. Analogy and Memory Anchor

Think of a cloud SOC as an international airport's security operation. At one level, physical screening machines check everyone against known prohibited items — the signature-based layer. At another level, behaviour analysts watch how people move through the terminal, flagging those whose behaviour deviates from normal passenger patterns — the behavioural analytics layer. At another level, specific individuals on watch lists trigger targeted alerts — the threat intelligence layer. At the top, a coordination centre correlates all of it and decides where to direct human response resources. The airport's security does not collapse because no single layer catches everything. It works because the layers complement each other, and trained humans make the final call on what to do.

Honeypot memory anchor: A honeypot is a tripwire with a digital trigger. No legitimate user or process should ever interact with it. If anything touches it, the alarm fires. The signal-to-noise ratio is near-perfect precisely because there is no legitimate reason for the interaction.

7. Exam Focus

ISC2 tests SOC concepts at the level of capability selection and architectural adequacy. Expect scenarios asking which monitoring capability would detect a specific type of threat, or evaluating whether a described SOC design adequately covers a given risk.

- Know the difference between IDS (detect and alert) and IPS (detect and block).
- Understand that honeypots generate high-fidelity alerts because any interaction is by definition suspicious.
- Know that behavioural analytics requires a baseline period before it is effective — a freshly deployed tool cannot detect deviations from normal if normal has not yet been established.

- Understand that intelligent monitoring reduces alert volume but does not eliminate human analyst requirements.
- Common distractor: 'Deploying AI-driven monitoring eliminates the need for a 24/7 SOC.' Incorrect — intelligent monitoring improves signal quality but human judgment for triage, escalation, and response remains essential.
- Know that cloud-native monitoring tools such as AWS GuardDuty and Azure Defender use ML against cloud-specific log sources (CloudTrail, VPC Flow Logs, DNS logs).

8. Key Terms Glossary

Security Operations Center (SOC): The organisational function responsible for continuous monitoring, detection, triage, and response to security threats.

Intrusion Detection System (IDS): A monitoring system that identifies suspicious traffic patterns and known attack signatures and generates alerts.

Intrusion Prevention System (IPS): A system that identifies and actively blocks suspicious traffic before it completes, rather than only alerting.

Honeypot: A deliberately exposed decoy system or resource with no legitimate business purpose, designed to attract and detect attacker interaction.

Network Security Group (NSG): A cloud platform control that filters traffic at the subnet or interface level and generates flow logs for analysis.

Signature-Based Detection: A monitoring approach that compares incoming data against a library of known attack patterns and generates alerts on matches.

Behavioural Analytics: A monitoring approach that establishes a baseline of normal activity and alerts on statistical deviations from that baseline.

Machine Learning (in SOC context): Automated pattern recognition applied to security event data to identify anomalies without requiring manual rule configuration for each threat type.

Alert Triage: The process by which a SOC analyst evaluates an incoming alert to determine its validity, severity, and appropriate response.

Runbook: A documented set of steps for responding to a specific type of security event, enabling consistent and rapid response under pressure.

AWS GuardDuty: A cloud-native threat detection service that uses ML to analyse CloudTrail, VPC Flow Logs, and DNS logs for anomalous or malicious activity.

Continuous Monitoring: Security monitoring that operates without interruption across all time periods, not limited to business hours.

9. Quick Revision Checklist

- ☐ I can describe the core mission of a SOC in one sentence.
- ☐ I can explain why continuous monitoring is a baseline requirement, not an advanced feature.
- ☐ I can name four distinct monitoring layers in a cloud SOC and what each detects.
- ☐ I can explain why honeypots generate high-fidelity alerts.
- ☐ I can differentiate IDS from IPS in terms of response action.
- ☐ I understand why signature-based detection alone is insufficient for modern threats.
- ☐ I can explain how behavioural analytics establishes a baseline and what it detects.
- ☐ I understand why intelligent monitoring does not eliminate the need for human analysts.
- ☐ I can name at least one cloud-native intelligent monitoring service and describe what data it analyses.
- ☐ I know what a runbook is and why it is essential to SOC operations.

10. Practice Questions

Scenario 1 (Questions 1–3): A cloud-hosted financial services company operates a SOC that runs 08:00–18:00 local time, Monday through Friday. Their monitoring stack includes signature-based IDS, firewall log analysis, and a SIEM that aggregates all log data. They have not yet deployed behavioural analytics or intelligent monitoring.

Q1. Which threat type is this SOC LEAST capable of detecting with its current configuration?

- A. A known ransomware variant that matches a published IDS signature.
- B. A brute-force login attack generating high-volume failed authentication events.
- C. A low-and-slow credential theft campaign that began on a Friday evening and operates only between 22:00 and 02:00.**
- D. A denial-of-service attack generating abnormal traffic volume.

Answer: C

A — Incorrect: Incorrect. Known ransomware variants that match published IDS signatures are exactly what signature-based detection is designed to catch.

B — Incorrect: Incorrect. A high-volume brute-force attack would generate significant log data and would likely trigger volume-based SIEM correlation rules.

C — Correct: Correct. A low-and-slow campaign operating outside SOC hours — after 18:00 Friday through Monday morning — has a 62-hour window of reduced human attention. Without behavioural analytics to detect deviations from normal patterns autonomously, and without 24/7 staffing, this campaign would proceed undetected.

D — Incorrect: Incorrect. A denial-of-service attack generating abnormal traffic volume would be visible in firewall logs and would likely trigger volume-based thresholds even without behavioural analytics.

Q2. The CISO proposes deploying a cloud-native behavioural analytics tool to address the detection gap. A team member raises the concern that the tool will have a high false positive rate. What is the MOST accurate characterisation of this concern?

A. The concern is unfounded; behavioural analytics tools do not generate false positives by design.

B. The concern is partially valid; newly deployed behavioural analytics tools require a baseline period and may generate elevated false positives until the baseline is established.

C. The concern is the primary reason behavioural analytics tools should not be deployed in financial services environments.

D. The concern is unfounded because the SIEM will automatically filter all behavioural analytics false positives before they reach analysts.

Answer: B

A — Incorrect: Incorrect. All detection systems can generate false positives. Behavioural analytics tools are no exception, particularly before a stable baseline is established.

B — Correct: Correct. Behavioural analytics tools must observe normal activity to establish a baseline. During the initial period — typically two to four weeks — the tool may flag legitimate activity as anomalous because it cannot yet distinguish normal from abnormal. This is a real and manageable limitation, not a reason to avoid deployment.

C — Incorrect: Incorrect. This overstates the limitation. False positive rates in established behavioural analytics tools are a tuning challenge, not a categorical disqualification for any industry.

D — Incorrect: Incorrect. SIEM correlation rules are not designed to automatically filter behavioural analytics false positives. That filtering requires human analyst review and rule tuning.

Q3. After deploying behavioural analytics, the SOC's on-call analyst receives an alert at 23:00 on a Friday. A service account that normally makes 12 API calls per night is making 340 calls across new endpoints. What is the CORRECT initial response?

- A. Disable the service account immediately and notify the CISO.
- B. Dismiss the alert as a false positive since no IDS signature was matched.
- C. Review the alert context to determine whether the activity represents a legitimate business change or a genuine anomaly before taking action.**
- D. Wait until Monday morning when the full SOC team is available to investigate.

Answer: C

A — Incorrect: Incorrect. Immediately disabling a service account without understanding the context risks disrupting a legitimate business process — for example, a month-end batch job running outside its normal window. Escalation should follow investigation, not precede it.

B — Incorrect: Incorrect. The absence of an IDS signature match does not make an anomaly a false positive. Behavioural analytics detects threats that signatures do not cover. Dismissing a behavioural alert because no signature fired defeats the purpose of deploying the tool.

C — Correct: Correct. The correct initial response is to evaluate the alert context. Is there a known change that would explain this behaviour? Has the application team deployed something new? Only after establishing that the activity cannot be explained by a legitimate change should the analyst escalate to containment actions.

D — Incorrect: Incorrect. Waiting until Monday morning gives a potential attacker the entire weekend to complete their operation. The on-call function exists precisely to address this scenario.

Q4. A SOC architect is designing the detection stack for a new cloud environment. Which statement BEST describes the relationship between signature-based IDS and behavioural analytics?

- A. Behavioural analytics replaces signature-based IDS because it detects a broader range of threats.
- B. Signature-based IDS replaces behavioural analytics in mature SOC environments because it generates fewer false positives.
- C. The two approaches are complementary: signature-based detection provides high-accuracy detection of known threats, while behavioural analytics extends coverage to novel and low-and-slow threats.**
- D. Behavioural analytics is only appropriate for environments with more than 10,000 users because smaller environments lack sufficient data for baseline establishment.

Answer: C

A — Incorrect: Incorrect. Behavioural analytics extends detection coverage but does not make signature-based detection redundant. Known threats covered by signatures should be caught with high accuracy and low latency — there is no benefit to waiting for a behavioural anomaly when a signature match is available.

B — Incorrect: Incorrect. Fewer false positives against known threats is a genuine advantage of signature-based detection, but it cannot compensate for the inability to detect novel threats. Maturity is not a reason to remove a capability layer.

C — Correct: Correct. The two approaches address different parts of the threat spectrum. Signature-based detection handles known threats efficiently. Behavioural analytics covers the blind spot for threats that signatures do not recognise. A mature SOC uses both in a layered architecture.

D — Incorrect: Incorrect. Behavioural analytics is beneficial in environments of all sizes. The baseline establishment process requires a meaningful period of normal operation, not a minimum user count.

Q5. Which of the following BEST explains why a honeypot generates higher-fidelity alerts than a production system monitored by IDS?

A. Honeypots use more sophisticated detection algorithms than standard IDS software.

B. Any interaction with a honeypot is inherently suspicious because no legitimate user or process has a reason to interact with it.

C. Honeypots generate encrypted alerts that cannot be tampered with by an attacker who has gained network access.

D. Honeypots monitor all network traffic in the environment, not just their own endpoints.

Answer: B

A — Incorrect: Incorrect. Honeypot fidelity does not come from more sophisticated detection algorithms. A simple honeypot with no automation at all generates high-fidelity alerts because of what triggers them, not how they are generated.

B — Correct: Correct. A honeypot's value is definitional: it has no legitimate purpose, so any interaction with it must be unauthorised. This eliminates the false positive problem that plagues production system monitoring, where distinguishing malicious from benign activity requires complex analysis.

C — Incorrect: Incorrect. Honeypot alerts are not distinguished by encryption. The fidelity advantage is from the nature of what triggers them, not how the alerts are transmitted.

D — Incorrect: Incorrect. A honeypot monitors only its own endpoints and interactions. It does not provide network-wide visibility.

Q6. Which cloud-native service uses machine learning to analyse CloudTrail API call logs, VPC flow logs, and DNS logs to detect threats without requiring manual rule configuration?

- A. AWS CloudWatch
- B. AWS Config
- C. AWS GuardDuty**
- D. AWS Inspector

Answer: C

A — Incorrect: Incorrect. CloudWatch is a monitoring and observability service for metrics, logs, and alarms. It does not apply machine learning to detect security threats autonomously.

B — Incorrect: Incorrect. AWS Config tracks configuration changes to AWS resources and evaluates them against compliance rules. It is not a threat detection service.

C — Correct: Correct. AWS GuardDuty is a cloud-native threat detection service that applies machine learning to CloudTrail, VPC flow logs, and DNS logs to identify malicious or anomalous behaviour without requiring manual signature or rule management.

D — Incorrect: Incorrect. AWS Inspector is a vulnerability assessment service that evaluates EC2 instances and container images for software vulnerabilities and network exposure. It does not analyse log data for threat behaviour.

Q7. A SOC analyst reviews an alert from the behavioural analytics platform indicating that a developer's account has logged in from a new geographic location at an unusual time. The analyst must decide whether to escalate or dismiss. What additional information is MOST relevant to the triage decision?

- A. Whether the developer's account has been active for more than 90 days.
- B. Whether the geographic location of the login corresponds to a known VPN exit node, and whether the developer recently notified IT of travel or remote work.**
- C. Whether the login generated an IDS signature match in addition to the behavioural alert.
- D. Whether the login occurred on a weekday or a weekend.

Answer: B

A — Incorrect: Incorrect. Account age is not relevant to the triage decision. A new or long-standing account can be equally compromised or legitimately used from a new location.

B — Correct: Correct. The most relevant contextual factors are whether the new location corresponds to a VPN (which could explain a location that looks foreign) and whether the developer has a known legitimate reason for the change. These two data points determine whether the anomaly is explainable or warrants escalation.

C — Incorrect: Incorrect. The absence of an IDS signature match does not increase or decrease the likelihood of compromise. A legitimate login from a new location also would not match an IDS signature.

D — Incorrect: Incorrect. While day of the week provides some context, it is far less relevant than understanding whether the deviation has a known legitimate explanation.

Scenario 2 (Questions 8–10): A healthcare organisation is designing its cloud SOC. The security architect presents two options: Option A deploys a signature-based IDS with a SIEM. Option B deploys a signature-based IDS, a SIEM, behavioural analytics, and honeypots in the development environment. The CFO is concerned about the additional cost of Option B.

Q8. The security architect argues that Option A is inadequate for the healthcare organisation's risk profile. Which threat scenario BEST supports this argument?

A. A known ransomware variant attacking unpatched systems.

B. An insider threat from a database administrator slowly exfiltrating patient records over six months by running queries slightly larger than normal.

C. A brute-force attack against the organisation's external login page.

D. A phishing email delivering a known malware payload with a published hash.

Answer: B

A — Incorrect: Incorrect. Known ransomware with an available signature is well within the detection capability of a signature-based IDS. This scenario does not support the argument for Option B.

B — Correct: Correct. A slow exfiltration by an insider — slightly larger queries, spread over months, by a user with legitimate access — will not match any IDS signature. The activity deviates from the administrator's established baseline behaviour, but only behavioural analytics would detect it. This is the exact gap Option A cannot close.

C — Incorrect: Incorrect. A brute-force attack generates high-volume failed login events that would trigger SIEM correlation rules. This scenario is detectable with Option A.

D — Incorrect: Incorrect. A phishing payload with a published hash would match IDS signatures. This scenario is within Option A's detection capability.

Q9. The CFO asks what specific value the honeypots in the development environment provide. What is the BEST response?

- A. Honeypots provide encryption for development environment traffic, reducing the risk of data interception.
- B. Honeypots act as decoy systems that attract attackers who have already penetrated the network perimeter, providing early warning of lateral movement.**
- C. Honeypots replace the need for network security group logging in the development environment.
- D. Honeypots provide automated remediation of detected threats, reducing analyst workload.

Answer: B

A — Incorrect: Incorrect. Honeypots do not provide traffic encryption. They are passive detection mechanisms, not network security controls.

B — Correct: Correct. Honeypots are particularly valuable for detecting lateral movement inside the network perimeter. An attacker who has breached the perimeter and is conducting reconnaissance may interact with a honeypot — a decoy system or resource that looks like a valid target. That interaction generates a high-fidelity alert that perimeter monitoring would not catch.

C — Incorrect: Incorrect. Honeypots do not replace NSG logging. They monitor only interaction with their own decoy endpoints and provide no visibility into broader network traffic.

D — Incorrect: Incorrect. Honeypots detect and alert; they do not remediate. Remediation is a human-led or automated response function separate from detection.

Q10. The CISO approves Option B. Six weeks after deployment, the SOC team reports that the behavioural analytics platform is generating a high volume of alerts on legitimate developer activity. What is the MOST appropriate response?

- A. Replace the behavioural analytics platform with a signature-based tool that generates fewer false positives.
- B. Disable the behavioural analytics alerts for the developer user group to reduce noise.
- C. Tune the platform's alert thresholds and exclusion rules based on six weeks of observed normal developer behaviour, and establish a feedback loop between analysts and the tuning process.**
- D. Accept the high alert volume as an inherent limitation of behavioural analytics and increase SOC headcount to process it.

Answer: C

A — Incorrect: Incorrect. The high alert volume six weeks into deployment is expected — the baseline is still being refined. Replacing the tool eliminates a capability that, once tuned, will provide significant detection value.

B — Incorrect: Incorrect. Disabling alerts for an entire user group creates a detection blind spot. A compromised developer account would go unmonitored.

C — Correct: Correct. Six weeks of observation provides meaningful behavioural data. The appropriate response is to use that data to tune the platform — adjusting thresholds, adding exclusion rules for known-normal developer workflows, and establishing a structured feedback process so analysts can flag specific false positives for rule adjustment.

D — Incorrect: Incorrect. Accepting high alert volume as permanent is operationally unsustainable and defeats the purpose of deploying the tool. Alert fatigue from untuned behavioural analytics is as dangerous as not having it — analysts begin to ignore or dismiss alerts.