

LEARNING RESOURCE GUIDE

ISC2 CCSP: Certified Cloud Security Professional (Aug 2026) | Professional Participant Workbook

Outsourcing and Cloud Contract Design: SLA, MSA, SOW

1. Learning Objectives

- By the end of this lecture, you will be able to Describe the function of the Master Service Agreement (MSA) and identify its key provisions in a cloud outsourcing relationship.
- By the end of this lecture, you will be able to Explain the purpose and content of a Statement of Work (SOW) and distinguish it from the MSA.
- By the end of this lecture, you will be able to Define a Service Level Agreement (SLA) and identify the metrics, measurement methodologies, and credit structures that a well-designed SLA should contain.
- By the end of this lecture, you will be able to Explain why an SLA is not a security guarantee and describe what additional contract instruments address security obligations.
- By the end of this lecture, you will be able to Evaluate the difference between a provider's standard contract terms and a negotiated enterprise agreement, and identify what is negotiable in practice.

2. Introduction: When the Platform Goes Dark

A global insurance broker runs its claims management platform on a hyperscale cloud service. On a Thursday morning in the middle of peak claims season, the platform goes offline. The cloud provider's status page shows a regional outage. Three hours pass. Then five. The broker's operations team logs hundreds of support requests. When service finally restores, the broker's management team convenes the post-incident review with one question at the centre: what did the contract say?

That question is not rhetorical. In cloud outsourcing relationships, the contract is the mechanism through which expectations are set, performance is measured, and accountability is enforced when things go wrong. A contract that promises broad availability commitments but contains a measurement methodology that excludes the exact failure mode that caused the outage is not protecting the customer — it is protecting the provider. Understanding the architecture of cloud contracts — what each document does, what each clause covers, and what is negotiable — is a core competency for cloud security professionals.

This lecture maps the three-document structure of cloud commercial relationships: the Master Service Agreement, the Statement of Work, and the Service Level Agreement. It clarifies what each instrument governs and what it does not, addresses the important distinction between SLA commitments and security guarantees, and examines the practical difference between standard provider terms and what can be achieved through enterprise negotiation.

3. Core Concepts Explained

3.1 Master Service Agreement (MSA) — The Legal Foundation

Definition: The Master Service Agreement is the foundational contract governing the entire commercial relationship between a customer and a service provider. It establishes the legal terms that apply across all engagements, regardless of the specific service or project.

Business relevance: The MSA is drafted to last. Individual projects and services come and go, but the MSA provides the legal framework that all of them operate within. For cloud relationships, it is the document that defines what happens when the relationship ends — including how data is returned or destroyed, what post-termination obligations exist, and under which jurisdiction disputes are resolved.

Key MSA provisions in cloud contracts:

Dispute resolution and governing law: The MSA specifies which country's law governs the contract and how disputes are resolved — litigation, arbitration, or alternative dispute resolution. For multinational organisations, this clause determines which legal system they are operating under for the entire cloud relationship.

Liability limitation: MSAs typically include caps on each party's maximum liability. Large cloud providers routinely cap their liability at a fraction of fees paid — often a few months' charges. This cap can be dramatically lower than the actual business loss from a major service failure, and it is one of the most important clauses to scrutinise and, where possible, negotiate.

Indemnification: Which party protects the other from third-party claims arising from specified events? In cloud contracts, indemnification clauses address scenarios such as a provider's intellectual property infringing a third party's rights, or a customer using the service in a way that generates third-party claims.

Data ownership and portability: The MSA should confirm that data the customer loads into the cloud environment remains the customer's property, and specify how that data will be returned or deleted upon contract termination. Weak provisions here create lock-in and complicate exit strategies.

Force majeure: Events outside either party's control that excuse non-performance. Cloud providers include force majeure clauses to exclude liability for outages caused by events such as natural disasters, government actions, or internet infrastructure failures. Customers in regulated industries should assess whether the force majeure definition is so broad that it effectively eliminates accountability for service continuity.

Confidentiality: How does each party protect the other's confidential information? For cloud customers, this covers proprietary business data, trade secrets, and — critically — the security configuration details shared with the provider for support purposes.

3.2 Statement of Work (SOW) — Defining the Engagement

Definition: A Statement of Work describes the specific deliverables, tasks, responsibilities, milestones, timelines, and acceptance criteria for a defined engagement. It operates under the MSA and fills in the specifics for a particular project or service.

Business relevance: The SOW is where scope is defined — and where scope disputes originate. If the SOW does not precisely define what constitutes a completed deliverable, who is responsible for which tasks, and how acceptance will be determined, both parties are exposed to disagreement about whether performance has been met.

In cloud contracts, SOWs are used for: migration projects (specifying what data moves, by when, and what success looks like), custom development or integration services, managed security services with defined outputs, and professional services engagements such as architecture design or compliance advisory work. A cloud relationship may have multiple active SOWs under a single MSA — each one scoped to a different piece of work, with its own timeline and acceptance criteria.

Key SOW elements: description of services and deliverables, milestone schedule with dates, responsibility matrix (who does what), acceptance criteria for each deliverable, payment schedule tied to milestone completion, and project governance arrangements including escalation paths.

Practitioner angle: The acceptance criteria clause in a SOW is the most frequently contested element. Without clear, measurable criteria, a vendor can argue that a deliverable meets the SOW's requirements even when the customer is genuinely dissatisfied. Acceptance criteria should be objective — defined in terms of measurable outputs, not subjective satisfaction.

3.3 Service Level Agreement (SLA) — Performance Commitments

Definition: A Service Level Agreement defines the measurable performance standards the provider commits to and the consequences — typically financial credits — when those standards are not met.

Business relevance: The SLA converts the general commitment to 'provide cloud services' into specific, verifiable performance obligations. It creates the mechanism for accountability when performance falls short and provides the customer with an objective basis for assessing whether the service is meeting business requirements.

Core SLA components:

Uptime commitments: The most familiar SLA metric. '99.9% monthly uptime' is a commonly advertised level, which equates to approximately 43.8 minutes of permitted downtime per month, or about 8.7 hours per year. '99.99%' equates to about 4.4 minutes per month. The number looks small until a failure at the wrong moment costs the customer significant revenue. Always assess uptime commitments against the actual business impact of outages during peak periods.

Response time objectives: How quickly will the provider respond to support incidents of different severity levels? A Severity 1 incident (complete service unavailability) might carry a 15-minute initial response commitment. A Severity 3 incident (minor degradation) might carry a 24-hour response time. These commitments should map to the customer's operational needs — a response time objective that does not align with business recovery requirements is not a meaningful commitment.

Recovery time objectives: SLAs sometimes include commitments on how long restoration will take after an outage begins. These are distinct from response commitments — response confirms the provider is aware and working; recovery specifies when service will be restored.

Measurement methodology: How is compliance calculated? Definitions matter enormously here. Does the provider's uptime calculation exclude scheduled maintenance windows? Does it exclude failures caused by the customer's own actions? A provider that measures uptime from their own monitoring infrastructure, excludes maintenance, and includes a broad 'customer-caused' exclusion may have a 99.95% uptime figure that masks actual availability levels the customer experiences.

Credit structures: What happens when the SLA is missed? Service credits are the standard remedy — typically a percentage of the monthly service fee for the affected service. AWS's Compute SLA offers 10% credit when monthly uptime falls below 99.9%, rising to 25% below 99.0% and 30% below 95.0%. Crucially, credits do not compensate for business losses — they offset a portion of the service fee. The gap between the credit value and the actual business impact of an outage is a risk the customer retains.

Claiming process: Credits are not automatic. The customer must file a claim within a defined window — often within 30 days of the incident — and the claim must be validated against the provider's measurement data. Failing to file within the window typically forfeits the credit entitlement.

Case Study: In March 2021, OVHCloud's SBG2 data centre in Strasbourg, France suffered a catastrophic fire that destroyed the facility and caused prolonged outages for customers across Europe. Multiple customers reported data loss, including organisations that had assumed the provider's infrastructure resilience was sufficient for their recovery requirements. Critically, many affected customers had not read the backup and disaster recovery provisions in their MSAs carefully. OVHCloud's standard terms had not guaranteed data backup for all service tiers. Post-incident analysis showed that customers who had separately contracted for managed backup services or who maintained their own offsite backups were able to recover. The incident reinforced a principle that applies to any cloud contract: the SLA governs service availability — it does not guarantee data protection unless data protection provisions are explicitly negotiated and contracted for.

3.4 SLA ≠ Security Guarantee

One of the most consequential misunderstandings in cloud governance is conflating the SLA with a security guarantee. An SLA is a performance instrument — it specifies availability, response times, and recovery commitments. It is not a contract about information security.

Security obligations come from different contract instruments: the MSA (which may include security requirements in its data protection annexe), a separately negotiated Data Processing Agreement (required for GDPR-regulated processing), and security addenda or schedules that specify encryption standards, access control requirements, incident response procedures, and audit rights. A customer who points to their SLA as evidence of their cloud security posture has confused availability with security — these are related but distinct concepts.

In practice, this means: read the security provisions in the MSA annexes, review the data processing agreement, check whether audit rights are included (the right to commission or receive the results of security audits), and ensure that incident notification timelines in the contract meet regulatory requirements — particularly the GDPR 72-hour notification obligation.

3.5 Standard Terms vs. Negotiated Enterprise Agreements

Large cloud providers publish standard terms of service. These are written by the provider's legal team for the provider's benefit. They typically: cap liability at a minimal percentage of fees, include broad force majeure exclusions, place the measurement burden entirely on the provider's monitoring data, and contain data protection provisions calibrated to minimum legal requirements rather than enterprise security standards.

What is negotiable in practice: Enterprise customers in regulated industries — financial services, healthcare, government — routinely negotiate modifications to standard terms. Common negotiated changes include: higher liability caps (especially for data breach scenarios), shorter breach notification timelines than the provider's standard (e.g., 24 hours rather than 72 hours), independent audit rights allowing the customer or their appointed auditor to assess the provider's controls, expanded data protection annexes meeting specific regulatory frameworks (FCA, HIPAA, PCI DSS, GDPR), and requirements for specific certifications to be maintained throughout the contract term.

The practical reality: the tier of service matters. A small business on a standard SaaS subscription has little negotiating leverage. An enterprise customer committing to seven or eight-figure annual spend typically does. Cloud security professionals should know what is worth negotiating, what the regulatory baseline requires, and when to escalate standard terms to legal counsel rather than accepting them as given.

Think of the contract structure as: a building lease arrangement. The MSA is the head lease — the overall legal relationship between landlord and tenant, including general obligations, liability framework, and termination conditions. The SOW is the fit-out schedule — specifying exactly what the landlord will deliver (office partitions,

IT cabling, lift access) before the tenant moves in, and by when. The SLA is the property management agreement — specifying standards for building maintenance, how quickly a broken lift will be repaired, and what rent discount applies if the heating fails for a week. Security obligations are a separate annex to the head lease — specifying CCTV coverage, access control standards, and incident response procedures for security events in the building.

4. Contract Document Comparison

Document	Function	What It Covers	What It Does NOT Cover	Duration
MSA	Legal foundation for the commercial relationship	Liability, indemnification, governing law, data ownership, confidentiality, termination, force majeure	Specific service performance levels; security control standards (covered in annexes)	Term of the overall relationship
SOW	Defines a specific engagement under the MSA	Deliverables, milestones, timelines, acceptance criteria, responsibilities, payment schedule	Ongoing service performance; legal relationship terms	Duration of the specific project or engagement
SLA	Defines measurable service performance commitments	Uptime, response times, credit structures, measurement methodology, claiming process	Security controls; data protection obligations; business loss compensation	Renewal cycle (typically annual or aligned to contract term)
DPA / Security Annex	Defines security and data protection obligations	Encryption standards, access controls, breach notification, audit rights, sub-processor controls	Service availability commitments; commercial terms	Co-terminous with MSA

5. SLA Uptime — What the Numbers Mean

SLA Level	Annual Permitted Downtime	Monthly Permitted Downtime	Suitability
99.0%	~87.6 hours / year	~7.3 hours / month	Non-critical workloads;

			development and test environments
99.9%	~8.7 hours / year	~43.8 minutes / month	Standard business applications; moderate criticality
99.95%	~4.4 hours / year	~21.9 minutes / month	Higher criticality; financial services standard tier
99.99%	~52 minutes / year	~4.4 minutes / month	High-availability requirements; payment processing, healthcare
99.999%	~5.3 minutes / year	~26 seconds / month	Mission-critical; telco-grade; rarely offered as a standard commitment

6. Exam Focus

- MSA = legal framework for the whole relationship: liability, indemnification, governing law, data ownership, termination.
- SOW = specific engagement details under the MSA: deliverables, milestones, acceptance criteria.
- SLA = measurable service performance commitments: uptime, response times, credits, measurement methodology.
- SLA credits compensate for service fee — not for business losses. The gap is retained by the customer.
- SLAs do not guarantee security. Security obligations are in the MSA annexes, DPA, and security schedules.
- Claims for SLA credits are not automatic. Filing deadlines apply.
- Standard provider terms are written for the provider's benefit. Enterprise customers in regulated industries negotiate modifications.
- Key negotiable terms: liability caps, breach notification timelines, audit rights, data protection annexes, certification maintenance requirements.
- Common distractor: treating an SLA as evidence of security compliance. SLAs cover availability, not security control effectiveness.
- Common distractor: assuming credits fully compensate for an outage. Credits are partial and limited to fee offsets.

7. Key Terms Glossary

Master Service Agreement (MSA): The foundational legal contract governing the overall commercial relationship between a customer and cloud provider, covering liability, indemnification, governing law, data ownership, and termination.

Statement of Work (SOW): A document under the MSA specifying the deliverables, responsibilities, timelines, milestones, and acceptance criteria for a specific engagement or project.

Service Level Agreement (SLA): A contract defining measurable service performance commitments — including uptime, response times, and credit structures — and the consequences of failing to meet them.

Uptime Commitment: The SLA metric specifying the percentage of time the service will be available within a defined measurement period, typically monthly.

Service Credit: A financial remedy applied against the customer's service fees when an SLA commitment is missed; calculated as a percentage of the monthly fee, not based on actual business loss.

Measurement Methodology: The defined process by which SLA compliance is calculated, including data sources, exclusions (e.g., scheduled maintenance, customer-caused failures), and measurement intervals.

Force Majeure: An MSA clause excusing non-performance due to events outside either party's control, such as natural disasters or government actions.

Liability Cap: An MSA provision limiting the maximum financial liability of either party, typically expressed as a multiple of monthly or annual fees.

Indemnification: An MSA clause under which one party agrees to protect the other from specified third-party claims.

Data Processing Agreement (DPA): A legally required contract between a GDPR data controller and processor specifying the terms, security requirements, and obligations for processing personal data.

Audit Right: A negotiated contract provision giving the customer the right to assess or commission assessment of the provider's controls, either directly or through the provider's audit reports.

Acceptance Criteria: The measurable, objective conditions specified in an SOW that define when a deliverable is complete and may be formally accepted by the customer.

8. Quick Revision Checklist

- ☐ I can describe the function of an MSA and name at least five key provisions it should contain.
- ☐ I can explain what a SOW covers and why precise acceptance criteria matter.
- ☐ I can define an SLA and list its core components: uptime, response times, measurement methodology, credit structures, claiming process.

- ☐ I understand that SLA credits offset service fees, not actual business losses.
- ☐ I can explain why an SLA does not constitute a security guarantee.
- ☐ I know that security obligations are governed by MSA annexes, DPAs, and security schedules — not the SLA.
- ☐ I understand that standard provider terms favour the provider and that enterprise customers routinely negotiate modifications.
- ☐ I can identify at least four terms that are commonly negotiated in enterprise cloud contracts.
- ☐ I can calculate the annual downtime permitted by a given SLA uptime percentage.
- ☐ I know that SLA service credit claims are not automatic and have filing deadlines.

9. Practice Questions (10 MCQs)

Question 1 (Scenario-Based Question)

Scenario: A regional energy company contracts with a cloud provider for infrastructure hosting services. Six months into the contract, the provider announces a significant infrastructure upgrade that will require planned maintenance windows totalling nine hours over the following month. The company's SLA specifies 99.9% monthly uptime. The provider's SLA measurement methodology excludes 'scheduled maintenance windows' from uptime calculations. After the maintenance month, the provider reports 99.95% uptime. What is the most accurate assessment of the company's SLA position?

A) The provider has breached the SLA because nine hours of downtime exceeds the 99.9% commitment's permitted downtime of 43.8 minutes. — Incorrect. Whether the SLA is breached depends on the measurement methodology in the contract. The SLA explicitly excludes scheduled maintenance from uptime calculations. The provider has reported 99.95% excluding maintenance — which satisfies the SLA as written.

B) The SLA has been met as defined by the contract, but the company should review whether the measurement methodology exclusions align with its actual availability requirements. — Correct. The SLA measurement methodology is contractual. The provider's reported figure of 99.95% excludes scheduled maintenance per the contract terms, so no technical breach exists. However, nine hours of planned downtime in a month represents a significant operational impact. The company should assess whether the exclusion clause is acceptable for its business needs and negotiate changes at renewal.

C) The SLA has been breached and the company is entitled to automatically receive service credits. — Incorrect. Based on the reported uptime figure of 99.95% excluding maintenance, there is no SLA breach. Additionally, even if a breach occurred, credits are not automatic — the customer must file a claim within the specified window.

D) The provider must compensate the company for all business losses incurred during the maintenance windows. — Incorrect. SLA credits are limited to a percentage of service fees. They do not compensate for business losses. The MSA liability cap further limits what the provider is obligated to pay.

Question 2 (Scenario-Based Question)

Scenario: A financial services firm is negotiating a cloud hosting contract with a hyperscale provider for its anti-money laundering (AML) processing platform. The provider's standard terms cap liability at three months' subscription fees, contain a 72-hour breach notification provision, and include no independent audit rights. The firm's compliance team identifies all three as inadequate for its regulatory obligations under the UK FCA's operational resilience rules. Which action best describes the appropriate procurement response?

A) Accept the standard terms and manage the compliance gap through internal controls. — Incorrect. The compliance team has identified specific provisions that do not meet regulatory requirements. Accepting non-compliant terms and attempting to compensate internally does not satisfy the FCA's operational resilience obligations, which require contractual assurance from third-party service providers.

B) Negotiate modifications to the MSA to increase the liability cap, shorten the breach notification window, and add contractual audit rights before signing. — Correct. Enterprise customers in regulated industries routinely negotiate MSA modifications. The liability cap, notification timeline, and audit rights are all negotiable in a significant enterprise contract. The firm's regulatory obligations provide clear grounds for requiring these modifications, and signing standard terms that conflict with regulatory requirements creates compliance exposure.

C) Request a SOC 2 Type II report from the provider instead, which would satisfy the FCA's requirements. — Incorrect. A SOC 2 Type II report provides assurance about control design and operating effectiveness but does not modify contractual terms. It does not address liability exposure, notification timelines, or audit rights — all of which are contractual matters requiring negotiation.

D) Proceed with signing under the standard terms, then rely on cyber insurance to cover the liability gap. — Incorrect. Cyber insurance addresses financial consequences after a loss event. It does not create compliance with FCA contractual requirements, does not give the firm real-time breach notification, and does not provide the audit visibility the regulator expects.

Question 3 (Scenario-Based Question)

Scenario: A professional services firm outsources its cloud-based document management platform to a SaaS provider. The project involves migrating 4 terabytes of client documents,

configuring role-based access controls, and integrating the platform with the firm's existing identity provider. The firm and provider have signed an MSA but the project team is now debating what constitutes successful completion of the migration. Which contract document should have prevented this dispute, and what specifically should it have contained?

A) The SLA; it should have specified response time commitments for the migration project. — Incorrect. SLAs govern ongoing service performance after go-live. They address uptime, response times, and credits — not project deliverable completion criteria.

B) The MSA; it should have included a detailed description of deliverables and timelines. — Incorrect. The MSA provides the legal framework for the overall relationship — governing law, liability, confidentiality. It is not the appropriate place for project-level delivery specifications.

C) The Statement of Work (SOW); it should have specified the deliverables, migration acceptance criteria, and responsibility matrix for each project phase. — Correct. The SOW is the instrument that defines project-specific deliverables, acceptance criteria, responsibilities, and timelines. A SOW specifying that migration is accepted when 100% of documents are successfully transferred and verified, all role-based access is configured and tested, and identity provider integration passes a defined functional test would have provided an objective basis for determining completion — and prevented the current dispute.

D) The Data Processing Agreement (DPA); it should have defined the migration timeline and data quality standards. — Incorrect. The DPA governs the legal basis and obligations for processing personal data. While relevant if the documents contain personal data, the DPA is not the instrument for defining project delivery specifications and acceptance criteria.

Question 4 (Knowledge Question)

Which of the following provisions is most commonly found in the Master Service Agreement (MSA) of a cloud contract, rather than in the SLA or SOW?

A) Uptime percentage and measurement methodology — Incorrect. Uptime commitments and measurement methodology are SLA provisions — they define performance standards for the service, not the overall legal relationship.

B) Project milestone dates and acceptance criteria for deliverables — Incorrect. Milestone dates and acceptance criteria are SOW provisions — they define the specific engagement, not the overarching legal framework.

C) Limitation of liability and governing law — Correct. Limitation of liability (capping each party's maximum exposure) and governing law (specifying which jurisdiction's law applies) are core MSA provisions. They define the legal parameters of the entire relationship, not any specific service performance commitment or project deliverable.

D) Service credit percentages applicable to availability failures — Incorrect. Service credit structures are SLA provisions — they define the financial consequences of specific performance failures.

Question 5 (Knowledge Question)

A cloud provider's SLA commits to 99.95% monthly uptime. The provider reports a month in which an unexpected outage caused 30 minutes of unavailability. Has the SLA been breached?

A) Yes, because 30 minutes of downtime exceeds the 99.9% threshold's allowance of 43.8 minutes. — Incorrect. The SLA in question is 99.95%, not 99.9%. The two thresholds have different permitted downtime allowances.

B) No, because 99.95% monthly uptime permits approximately 21.9 minutes of downtime per month, and 30 minutes of unavailability exceeds that threshold. — Incorrect. This answer contains the right calculation but draws the wrong conclusion. 30 minutes > 21.9 minutes means the SLA HAS been breached — not that it has not.

C) Yes, because 30 minutes of unavailability exceeds the 99.95% uptime commitment's permitted downtime of approximately 21.9 minutes per month. — Correct. 99.95% monthly uptime permits approximately 21.9 minutes of downtime in a 30-day month ($720 \text{ hours} \times 0.0005 = 0.36 \text{ hours} = 21.6 \text{ minutes}$). 30 minutes of unavailability exceeds this threshold, constituting a breach. The customer should file a credit claim within the contract's stipulated window.

D) Impossible to determine without knowing whether the outage was excluded by the measurement methodology. — Partially valid in practice, but the question presents a straightforward arithmetic scenario. Based on the figures given, 30 minutes > 21.9 minutes = breach. In a real contract, measurement methodology exclusions would be the next thing to check.

Question 6 (Knowledge Question)

AWS's published Compute SLA offers service credits of 10% when monthly uptime drops below 99.9%, and 30% when it drops below 95.0%. A customer's monthly service fee is \$50,000 and they experienced an outage that caused monthly uptime of 93%. What is the maximum credit the customer can claim?

A) \$5,000 (10% of monthly fee) — Incorrect. 93% uptime falls below the 95% threshold, which triggers the 30% credit tier, not the 10% tier.

B) \$15,000 (30% of monthly fee) — Correct. 93% monthly uptime is below the 95% threshold that triggers the 30% credit rate. $30\% \times \$50,000 = \$15,000$. Note that this credit

offsets a portion of the monthly service fee — it does not compensate for actual business losses from the outage.

C) \$50,000 (full monthly fee refund) — Incorrect. SLA credits are percentage-based offsets against the service fee, not full refunds. Cloud providers do not typically offer 100% credit under standard SLA structures.

D) The full business loss incurred during the outage period — Incorrect. SLA credits are explicitly limited to a percentage of service fees and are not designed to compensate for business losses. The gap between the credit value and actual business impact is a risk retained by the customer.

Question 7 (Knowledge Question)

A customer's cloud contract includes a strong SLA committing to 99.99% monthly uptime with meaningful credit structures. Following a security incident in which the provider's management plane was briefly accessed by an unauthorised party, the customer seeks to enforce remedies under the SLA. What limitation will the customer encounter?

A) SLA credits can only be claimed in the quarter following the incident. — Incorrect. SLA credit claim windows are defined in the contract, but timing restrictions are not the fundamental limitation here.

B) The SLA governs service availability, not security incidents. Remedies for security breaches are addressed through separate contract provisions — the MSA, security annexes, and DPA — not the SLA. — Correct. SLAs are performance instruments — they address uptime and availability. A security incident that does not cause a service outage may not trigger any SLA remedy at all. Security incident obligations, notification requirements, and remedies are governed by other contract instruments. The customer's security-related remedies depend on what was negotiated in the MSA and security annexes.

C) The customer cannot claim any remedy because the incident involved the management plane rather than the customer-facing service. — Incorrect. While the SLA may not be triggered if service was uninterrupted, the customer may have remedies under other contract provisions — particularly if security incident notification or breach obligations were violated.

D) SLA credits will automatically be issued at the maximum rate because a security incident constitutes a force majeure event. — Incorrect. Force majeure clauses typically excuse performance obligations — they do not trigger automatic maximum credits. A security incident would not ordinarily qualify as force majeure, and SLA credits are not automatic regardless.

Question 8 (Knowledge Question)

Which of the following is the most accurate description of a force majeure clause in a cloud MSA?

A) A provision requiring the provider to notify the customer within 24 hours of any event affecting service availability. — Incorrect. Notification requirements are typically in the SLA or security annexes. Force majeure clauses address excuse of performance, not notification obligations.

B) A clause excusing either party from performance obligations arising from events outside their reasonable control, such as natural disasters or government actions. — Correct. Force majeure clauses identify categories of unforeseeable events beyond either party's control that excuse non-performance. Cloud customers should scrutinise the breadth of the force majeure definition — an overly broad clause can effectively immunise the provider from accountability for events that, with proper resilience planning, could have been managed.

C) A limitation on the maximum service credits the customer can claim in any single calendar year. — Incorrect. Annual credit caps may exist in SLAs, but they are distinct from force majeure clauses, which address excuse of performance obligations rather than limiting financial remedies.

D) A provision requiring the provider to maintain geographically distributed infrastructure as a condition of service delivery. — Incorrect. Infrastructure resilience requirements are negotiated separately, often in the technical specifications or service description. Force majeure addresses events that excuse performance, not the baseline infrastructure obligations.

Question 9 (Knowledge Question)

Why are standard cloud provider contract terms generally less favourable to enterprise customers in regulated industries than negotiated terms?

A) Standard terms are written by regulators and must be accepted without modification. — Incorrect. Standard terms are written by the cloud provider's legal team, not regulators. Regulators may prescribe certain minimum standards that contracts must meet, but the contract itself is a commercial instrument authored by the provider.

B) Standard terms are drafted by the provider to minimise the provider's liability and exposure, and do not account for the specific regulatory obligations or risk profile of each enterprise customer. — Correct. Standard terms are written at scale for the provider's commercial benefit. They typically cap liability at minimal levels, include broad exclusions, and contain data protection provisions calibrated to minimum legal requirements. Enterprise customers in regulated industries — who face specific regulatory obligations around data protection, operational resilience, and audit rights — routinely negotiate modifications to standard terms.

C) Standard terms are updated annually and always reflect current best practice. — Incorrect. Standard terms are updated at the provider's discretion and are designed to protect the provider's interests. Currency does not equate to alignment with any particular customer's regulatory requirements.

D) Standard terms cannot be modified once the contract is signed. — Incorrect. While modifications after signing are more difficult, the point here is that negotiation occurs before signing. Enterprise customers negotiate modifications prior to contract execution.

Question 10 (Knowledge Question)

A cloud customer experiences a service outage that breaches the provider's SLA uptime commitment. The outage results in \$800,000 of lost business revenue. The customer's monthly service fee is \$30,000 and the applicable credit rate is 25%. What is the maximum SLA credit the customer can claim, and what happens to the remaining loss?

A) The customer can claim \$800,000 because the SLA compensates for actual business losses. — Incorrect. SLAs do not compensate for actual business losses. Credits are calculated as a percentage of service fees — not based on the financial impact of the outage.

B) The customer can claim \$7,500 (25% of the \$30,000 monthly fee). The remaining \$792,500 of business loss is retained by the customer unless separately recovered through the MSA liability framework or cyber insurance. — Correct. $\text{SLA credit} = 25\% \times \$30,000 = \$7,500$. This offsets a portion of the monthly fee. The \$792,500 gap between the credit and the actual loss is the customer's retained risk — illustrating why SLA credits are not designed to compensate for business impact. Recovery of further losses depends on the MSA liability provisions (if the outage constitutes a breach of the MSA) and any applicable insurance coverage.

C) The customer can claim \$200,000 (25% of the business loss). — Incorrect. Credits are applied to the service fee, not the business loss. The 25% rate is applied to \$30,000 (the monthly fee), not to \$800,000.

D) No credit is available because the outage was a single event rather than a pattern of recurring failures. — Incorrect. A single event that breaches the uptime threshold in a given measurement period is sufficient to trigger a credit claim. There is no requirement for repeated failures.